

Last updated: **Tuesday 16th October 9:52pm**
 Most recent changes are shown in **red**;
 older changes are shown in **brown**.

Aims

This assignment aims to give you an understanding of

- how database files are structured and accessed
- how superimposed codeword (SIMC) signatures are implemented
- how partial-match retrieval searching is implemented using SIMC signatures

The goal is to build a simple implementation of a SIMC signature file, including application to create SIMC files, insert tuples into them, and search for tuples based on partial-match retrieval queries.

Summary

Deadline: 23:59:59pm on **Sunday 21 October**

Late Penalty: 0.09 *marks* off the ceiling mark for each hour late

Marks: Contributes **15 marks** toward your total mark for this course.

Groups: do this assignment in pairs or individually (you can use the same groups as for Assignment 1)

Submission: Login to Course Web Site > Assignments > Assignment 2 > Submission > upload `ass2.tar`

The `ass2.tar` file must contain the `Makefile` plus all of the `*.c` and `*.h` files that are needed to compile the `create`, `insert` and `select` executables. However, you should not change or submit `create.c`, `insert.c` and `select.c`. Details on how to build the `ass2.tar` file are given below.

Make sure that you read this assignment specification *carefully and completely* before starting work on the assignment. Questions which indicate that you

Note: this assignment does not r

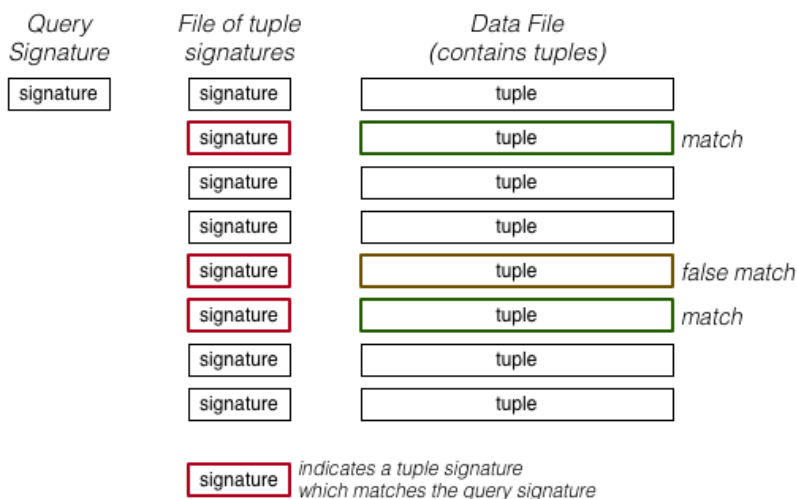
<https://eduassistpro.github.io/>

Introduction

Signatures are a style of indexing where (in its simplest form) each tuple They are used in the context of partial-match retrieval queries, and are p scanning signatures, matching them against a query signature, and then Efficient signature matching (small signatures, simple bit-comparison) all match, but the tuple is not a valid result for the query.

representation of its values. des. Selection is performed by ged as potential matches. e the query and tuple signatures

The kind of signature matching described above uses one signature for each tuple (as in the diagram below). Other kinds of signatures exist, and one goal is to implement them and compare their performance to that of tuple signatures.



Signatures can be formed in several ways, but we will consider only signatures that are formed by superimposing codewords (SIMC). Each codeword is formed using the value in one attribute.

| | |
|------------------|------------------|
| codeword(1) | 0100001000100000 |
| codeword(a) | 1000000100000010 |
| codeword(42) | 0010000000000010 |
| signature(1a,42) | 1110001100100010 |

Forming a signature
Tuple = (1,a,42)

In our context, SIMC-indexed relations are a collection of files that represent one relational table, and can be manipulated by a number of supplied commands:

gendata #tuples #attributes [startID] [seed]

Generates a specified number of n -attribute tuples in the appropriate format to insert into a created relation. All tuples are the same format and look like

UniqID,RandomString,a3-Num,a4-Num,...,a_n-Num

For example, the following 4-attribute tuples could be generated by a call like `gendata 1000 4`

7654321,aTwentyCharLongStrng,a3-013,a4-001
 3456789,aTwentyChrLongString,a3-042,a4-128

A tuple is a sequence of comma-separated fields. The first field is a unique 7-digit number; the second field is a random 20-char string; the remaining fields have a field identifier followed by a non-unique 3-digit number. The size of each tuple is

$7+1 + 20+1 + (n-2)*(6+1) - 1 = 28 + 7*(n-2)$ bytes

Note that tuples are limited to at most 9 attributes, which means that the maximum tuple size is a modest 77 bytes. (If you wish, you can work with larger tuples by tweaking the `gendata` and `create` commands and the `newRelation()` function, but this not required for the assignment)

create RelName #tuples #attrs 1/pF

Creates an empty relation number of tuples that can probability; for example, a

les parameter give a maximum
 ves the inverse of the false match

These parameters are combined using the formulas given in lectures to determine how large tuple- and page-signatures are. Each bit-slice has a number of bits equal to the number of data pa size. (Aside: this is different to what I said in lectures, where I was thinking to especially if the relation is approaching full); this is slightly more realistic,

- the name of the relation
- the number of attributes
- the false match probability

This gives you storage for one relation/table, and is analogous to making an SQL data definition like:

`create table R (a1 integer, a2 text, ... an text);`

Note that internally, attributes are indexed 0.. $n-1$ rather than 1.. n .

The following example of using `create` makes a relation called `abc` where each tuple has 4 attributes and the indexing has a false match probability of 1/100. The relation can hold up to 10000 tuples (it can actually hold more, but only the first 10000 will be indexed via the bit-sliced signatures).

`$./create abc 10000 4 100`

insert RelName

Reads tuples, one per line, from standard input and inserts them into the relation specified on the command line. Tuples all take the form *val₁,val₂,...,val_n*. The values can be any sequence of alpha-numeric characters and '-'. The characters ',' (field separator) and '?' (query wildcard) are treated specially.

Since all tuples need to be the same length, it is simplest to use `gendata` to generate them, and pipe the generated tuples into the `insert` command

dump RelName

Writes all tuples from the relation *RelName*, one per line, to standard output. This is like an inverse of the `insert` command. Tuples are dumped in a form that could be used by `insert` to rebuild a database.

select RelName QueryString SigType

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Takes a "query tuple" on the command line, and finds all tuples in the data pages of the relation *RelName* that match the query. Queries take the form $val_1, val_2, \dots, val_n$, where some of the val_i can be '?' (without the quotes). Some examples, and their interpretation are given below. You can find more examples in the lecture slides and course notes.

```
? , ? , ?      # matches any tuple in the relation
10 , ? , ?     # matches any tuple with 10 as the value of attribute 1
? , abc , ?    # matches any tuple with abc as the value of attribute 2
10 , abc , ?   # matches any tuple with 10 and abc as the values of attributes 1 and 2
```

Note that the `startQuery()` function can return NULL. It should do so only if the query string contains the wrong number of attributes for the relation.

In fact, a SIMC relation *R* is represented by five physical files:

- *R.info* containing global information such as
 - the number of attributes and size of each tuple
 - the number of data pages and number of tuples
 - the sizes of the various kinds of signatures
 - the number of signatures and signature pages
 - etc. etc. etc.

The *R.info* file contains a copy of the *RelnParams* structure given in the *reln.h* file (see below).

- *R.data* containing data pages, where each data page contains
 - a count of the number of tuples in the page
 - the tuples (as comma-separated character sequences)

Each data page has a capacity of *c* tuples. If there are *n* tuples then there will be $b = \lceil n/c \rceil$ pages in the data file. All pages except the last are full. Tuples are never deleted.

- *R.tsig* containing tuple signatures, where each page contains
 - a count of the number of signatures in the page
 - the signatures themselves (as bit strings)

Each tuple signature is for

tuple. If there are *n* tuples in the relation, there will be *n* tuples in the *R.tsig* file.

- *R.psig* containing page signatures
 - a count of the number of signatures in the page
 - the signatures themselves (as bit strings)

Page signatures are much larger than tuple signatures, and are for all tuples in the page. There is one page signature for each page.

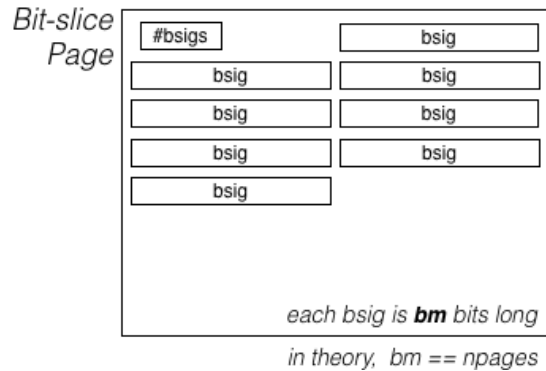
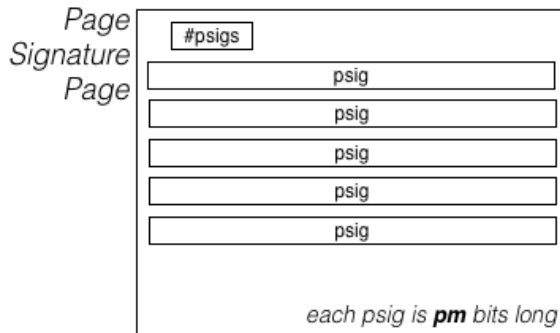
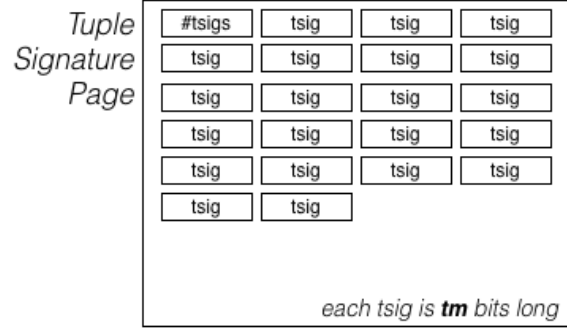

- *R.bsig* containing bit-sliced signatures, where each page contains
 - a count of the number of signatures in the page
 - the bit-slices themselves (as bit strings)

Bit-slices give an alternate 90° view of page signatures. If there are *b* data pages, then each bit-slice is *b*-bits long. If page signatures are *pm* bits long, then there are *pm* bit-slices. attribute values in all tuples in the page.

The following diagram gives a very simple example of the correspondence between page signatures and bit-slices:



The different types of pages (tuple, signature, slice) were described above. Internally, all pages have a similar structure: a counter holding the number of items in the page, and the items themselves (tuples or signatures or slices). All of the items in a page are the same size. The following diagram shows the structure of pages in the files of a SIMC relation:



We have developed some infrastructure for you to use in implementing SIMC files. You may use this infrastructure or replace parts of it (or all of it) with your own, but your SIMC files implementation must conform to the conventions used in our code. In particular, you should preserve the interfaces to the supplied modules (e.g. Reln, Query, Tuple) and ensure that your submitted modules work with the supplied code in the create, insert, and select commands.

Assignment Project Exam Help

Setting Up

You should make a working directory <https://eduassistpro.github.io/> if reading to make sure that you understand all of the data types

```
$ mkdir your/ass2/directory
$ cd your/ass2/directory
$ unzip /web/cs9315/18s1/assignments/ass2/ass2.zip
```

Add WeChat edu_assist_pro

You should see the following files in the directory:

```
$ ls
Makefile      dump.c        psig.c        stats.c       x1.c
bits.c        gendata.c    psig.h        tsig.c        x2.c
bits.h        hash.c       query.c       tsig.h        x3.c
bsig.c        hash.h       query.h       tuple.c
bsig.h        insert.c     reln.c        tuple.h
create.c      page.c      reln.h        util.c
defs.h        page.h      select.c      util.h
```

The .h files define data types and function interfaces for the various types used in the system. The corresponding .c files contain the implementation of the functions on the data type. The remaining .c files either provide the commands described above, or are test harnesses for individual types (x1.c, x2.c, x3.c). You can add additional testing files, but there is no need to submit them.

The above files give you a partial implementation of SIMC indexing. You need to complete the code so that it provides the functionality described above.

You should be able to build the supplied partial implementation via the following:

```
$ make
gcc -std=gnu99 -Wall -Werror -g -c -o query.o query.c
gcc -std=gnu99 -Wall -Werror -g -c -o page.o page.c
gcc -std=gnu99 -Wall -Werror -g -c -o reln.o reln.c
gcc -std=gnu99 -Wall -Werror -g -c -o tuple.o tuple.c
gcc -std=gnu99 -Wall -Werror -g -c -o util.o util.c
gcc -std=gnu99 -Wall -Werror -g -c -o tsig.o tsig.c
gcc -std=gnu99 -Wall -Werror -g -c -o psig.o psig.c
gcc -std=gnu99 -Wall -Werror -g -c -o bsig.o bsig.c
gcc -std=gnu99 -Wall -Werror -g -c -o hash.o hash.c
gcc -std=gnu99 -Wall -Werror -g -c -o bits.o bits.c
gcc -std=gnu99 -Wall -Werror -g -c -o create.o create.c
gcc -o create create.o query.o page.o reln.o tuple.o util.o tsig.o psig.o bsig.o hash.o bits.o -lm
```

```
gcc -std=gnu99 -Wall -Werror -g -c -o insert.o insert.c
gcc insert.o query.o page.o reln.o tuple.o util.o tsig.o psig.o bsig.o hash.o bits.o -o insert
gcc -std=gnu99 -Wall -Werror -g -c -o select.o select.c
gcc select.o query.o page.o reln.o tuple.o util.o tsig.o psig.o bsig.o hash.o bits.o -o select
gcc -std=gnu99 -Wall -Werror -g -c -o stats.o stats.c
gcc stats.o query.o page.o reln.o tuple.o util.o tsig.o psig.o bsig.o hash.o bits.o -o stats
gcc -std=gnu99 -Wall -Werror -g -c -o gendata.o gendata.c
gcc -o gendata gendata.o util.o -lm
gcc -std=gnu99 -Wall -Werror -g -c -o dump.o dump.c
gcc dump.o query.o page.o reln.o tuple.o util.o tsig.o psig.o bsig.o hash.o bits.o -o dump
gcc -std=gnu99 -Wall -Werror -g -c -o x1.o x1.c
gcc -o x1 x1.o query.o page.o reln.o tuple.o util.o tsig.o psig.o bsig.o hash.o bits.o
gcc -std=gnu99 -Wall -Werror -g -c -o x2.o x2.c
gcc -o x2 x2.o query.o page.o reln.o tuple.o util.o tsig.o psig.o bsig.o hash.o bits.o
gcc -std=gnu99 -Wall -Werror -g -c -o x3.o x3.c
gcc -o x3 x3.o query.o page.o reln.o tuple.o util.o tsig.o psig.o bsig.o hash.o bits.o
```

This should not produce any errors on the CSE servers; let me know ASAP if this is not the case.

The gendata command should work completely without change. For example, the following command generates 5 tuples, each of which has 4 attributes. Values in the first attribute are unique; values in the second attribute are highly likely to be unique. Note that the third and fourth attributes cycle through values at different rates, so they won't always have the same number.

```
$ ./gendata 5 4
1000000,lrfkQyuQFjKXyQVNRTyS,a3-000,a4-000
1000001,FrzrmzLYGFvEulQfpDBH,a3-001,a4-001
1000002,lqDqrrCRwDnXeuOQqekl,a3-002,a4-002
1000003,AITGDPHCSPiJtHbsFyfv,a3-003,a4-003
1000004,lADzPBfudkKlrwqAOzMi,a3-004,a4-004
```

The create command itself is complete, but some of the functions it calls are not complete. It will allow you to make an empty relation, although without a complete bit-slice file (you add this as one of the assignment tasks). The stats command is complete and can display information about a relation. Using these commands you could do the following: use the create command to create an empty relation which can hold 4 attribute tuples and able to index up to 5000 tuples (using bit slices), with a false match probability of 1/1000. The stats command then displays the generated parameter values.

```
$ ./create R 5000 4 1
$ ./stats R
Global Info:
Dynamic:
#items: tuples: 0 tsigs: 0 psigs: 0 bsigs: 0
#pages: tuples: 1 tsigs: 1 psigs: 1 bsigs: 1
Static:
tups #attrs: 4 size: 42 bytes max/page: 97
sigs bits/attr: 9
tsigs size: 64 bits (8 bytes) max/page: 511
psigs size: 5584 bits (698 bytes) max/page: 5
bsigs size: 56 bits (7 bytes) max/page: 584
$
```

<https://eduassistpro.github.io/>
Add WeChat edu_assist_pro

You can apply the formulae for calculating the various quantities to check that the above values make sense. Note that the bits for signatures are rounded up to the next multiple of 8 (why waste a few bits?). Note also that all pages are defined to be 4096 bytes. Finally, note that create makes a file with one empty page for each of the files holding tuples, signatures and slices.

As supplied, the insert command inserts tuples into the data pages, but does not generate any signatures. Using gendata is the easiest (and safest) way to add valid tuples. You can then check that the tuples have been inserted via the dump command, and see how the parameters have changed using stats again.

```
$ ./gendata 5 4 | ./insert R
Inserting: 1000000,lrfkQyuQFjKXyQVNRTyS,a3-000,a4-000
Inserting: 1000001,FrzrmzLYGFvEulQfpDBH,a3-001,a4-001
Inserting: 1000002,lqDqrrCRwDnXeuOQqekl,a3-002,a4-002
Inserting: 1000003,AITGDPHCSPiJtHbsFyfv,a3-003,a4-003
Inserting: 1000004,lADzPBfudkKlrwqAOzMi,a3-004,a4-004
$ ./stats R
Global Info:
Dynamic:
#items: tuples: 5 tsigs: 0 psigs: 0 bsigs: 0
#pages: tuples: 1 tsigs: 1 psigs: 1 bsigs: 1
Static:
tups #attrs: 4 size: 42 bytes max/page: 97
sigs bits/attr: 9
tsigs size: 64 bits (8 bytes) max/page: 512
psigs size: 5584 bits (698 bytes) max/page: 5
bsigs size: 56 bits (7 bytes) max/page: 585
$
```

Note that the only difference between the above stats and the stats for the newly-created file is the 5 tuples. There are no signatures, no new pages, etc.

The `dump` command is complete; it simply scans the data file and displays any tuples it finds, e.g.

```
$ ./dump R
1000000,lrfkQyuQFjKXyQVNRTyS,a3-000,a4-000
1000001,FrzrmzLYGFvEulQfpDBH,a3-001,a4-001
1000002,lqDqrrCRwDnXeuOQgekl,a3-002,a4-002
1000003,AITGDPHCSPiJtHbsFyfv,a3-003,a4-003
1000004,lADzPBfudkKlrwqAOzMi,a3-004,a4-004
$
```

The `select` command, as supplied, is not complete. However, once it is working (at least with tuple signatures), you should be able to ask queries like:

```
$ ./select R 1000001,?,? t      # not enough attrs
Invalid query: 101,?,?
$ ./select R 1000001,?,?,? t
Matched Pages:00000001
1000001,FrzrmzLYGFvEulQfpDBH,a3-001,a4-001
Query Stats:
# signatures read: 5
# sig pages read: 1
# tuples examined: 5
# data pages read: 1
# false match pages: 0
$ ./select R 1000001,?,a3-002,? t
Matched Pages:00000000
Query Stats:
# signatures read: 5
# sig pages read: 1
# tuples examined: 0
# data pages read: 0
# false match pages: 0
$ ./select R 1000001,?,a3-002,? x
Query Stats:
# signatures read:
# sig pages read:
# tuples examined:
# data pages read: 1
# false match pages: 1
```

Some explanation:

- The second query finds a match because there is a tuple with the value 1000001 for its first attribute. The ? represent "don't care" or wild-card values.
- The third query fails because the tuple with 1000001 for its first attribute, does not have the value a3-002 for its third attribute.
- The fourth query performs a linear scan of the data file. Since the query itself is the same as the third query, there are no matching tuples. This query reads every data page (there is only one). Any data page read, which does not contain matching tuples, is counted as a "false page match".
- The digits after `Matched Pages` is a bit-string, with one bit for each data page. The bit-string contains a 1 for each page where the indexing thinks that there is a match. In this case, there is only one data page, so there's only one bit that could be set. There are 8 bits in the "matched pages" string, because we round the size of all bit-strings up to a multiple of 8.
- The `t` at the end of the query tells the query evaluator to use tuple signatures as a first-pass filter. Other possibilities are `p` for page signatures or `b` for bit-sliced signatures. If you don't specify a signature type, the evaluator uses a liner scan and checks all tuples.
- With all types of signatures, queries run in two phases:
 - use the signatures to determine pages which potentially contain matching tuples
 - read each of these pages and check all tuples to find real matches
- The query statistics are maintained in a `Query` data structure while the query is executing.

Data Types


There are four important data types defined in the system:

Relations (data type `Reln`)

Relations are defined by three data types: `Reln`, `RelnRep`, `RelnParams`. `Reln` is just a pointer to a `RelnRep` object; this is useful for passing to functions that need to modify some aspect of the relation structure. `RelnRep` is a representation of an open relation and contains the parameters, plus file descriptors for all of the open files. `RelnParams` is a list of various properties of the database. See `reln.h` for details.

Queries (data type Query)

Queries are defined via a `QueryRep` structure which contains fields to represent the current state of the scan for the query, plus a collection of statistics counters. It is essentially like the query iteration structures described in lectures, and is used to control and monitor the query evaluation. The `QueryRep` structure also contains a reference to the relation being queried, and a copy of the query string. The `Query` data type is simply a pointer to a `QueryRep` structure. See `query.h` for details. The following diagram might also help:



Pages (data type Page)


Pages are defined via a `PageRep` structure which contains a counter for the number of items, and then an array of bytes containing the actual items, whether they are tuples or signatures or slices. The size of each type of item is held in the `RelnParams` structure, and so Pages are typically considered in conjunction with `Relns`. The `Page` data type is simply a pointer to a `PageRep` structure. See `page.h` for details. The following diagram might also help:

Assignment Project Exam Help

Bit-strings (data type Bits)

Bit-strings are defined via a `BitsRep` structure which contains a counter for the number of bytes used to represent the bit-string. The `BitsRep` structure is hidden from its clients; the `Bits` data type is simply an ADT, so the concrete structure. See `bits.c` for details.

number of bits, and the other for the number of bytes which hold the bits in the `BitsRep` structure. See `bits.c` for details.



Tuple (data type Tuple)

Tuples are just character sequences (like C strings). See `tuple.h` for details.

There are also a range of (hopefully) self-explanatory data types defined in `defs.h`. The various signature types are represented as bit-strings (`Bits`).

Goal

Your goal for this assignment is to complete the implementation of the various components of the system, so that it can handle all three kinds of signatures. This includes inserting/updating signatures when new tuples are added, and using these signatures in answering queries.

The header (`.h`) files contain definitions of the data types used in the system.

Each of the source code (`.c`) files contains comments on each function, describing briefly what it should do. Some functions contain `//TODO` comments to indicate where you need to complete them. You can put all of the code in the indicated function, or you can write new functions that these functions use.

You are free to change any file except `create.c`, `insert.c` and `select.c`. Since you can't change these files, you also cannot change the interfaces to the data types that they use (`Reln`, `Query`, `Page`, `Bits`). Basically, all of the functions mentioned in the `.h` files for these types must exist, with the same interface, but you can implement their internals however you like. You can also add extra functions to each data type (i.e. extend its interface) if that helps.

Task 1: A Bit-string Type (2 marks)

Implement all of the incomplete functions in the `bits.c` file, to produce a working bit-string data type. The functions to complete are flagged with `TODO`, and the purpose of each should be clear from the comment at the start of the function and its name. The `x1.c` file contains some simple test cases for the `Bits` type.

Task 2: Scanning for Results (2 marks)

After you have `Bits` working, you can start to implement query evaluation, although without indexing. The `startQuery()` function parses the query string and then uses the appropriate type of signature to generate a list of pages which potentially contain matching tuples. This list is implemented as a bit-string where a 1 indicates a page which needs to be checked for matches. At this stage, all of the signature types mark all pages as potential matches, so all pages need to be checked.

Before this will work, you need to implement the `scanAndDisplayMatchingTuples()` function, which performs the check for matching tuples in each of the marked pages. This function, as well as finding and displaying result tuples, maintains the query statistics for number of data pages read, and number of pages that were read but contained no matching tuples.

For this task, you need to complete the `scanAndDisplayMatchingTuples()` function from the `query.c` file. This function behaves roughly as follows:

```
foreach PID in 0 .. npages-1 {
    if (PID is not set in MatchingPages)
        ignore this page
    for each tuple T in page PID {
        if (T matches the query string)
            display it as a query result
    }
    if (no tuples in page PID are results)
        count it as a false match page
}
```

Task 3: Tuple Signatures (4 marks)

Implement indexing by using tuple-based signatures (i.e. each tuple has its own signature, stored in the `Rel.tsig` file). You will need to complete the `makeTupleSig` and `addToRelation()` function in the `rel.c` file, and add some code to the `rel.h` file.

The `addToRelation()` function currently does nothing with signatures. You should add code here which adds a tuple to the next available slot in the `Rel.tsig` file.

The `makeTupleSig()` function takes a tuple and returns a bit-string which is the codeword signature for that tuple. It behaves roughly as follows:

```
Tsig = AllZeroBits
for each attribute A in tuple T {
    CW = codeword for A
    Tsig = Tsig OR CW
}
```

A method for computing codewords is given in the lecture notes.

The `findPagesUsingTupSigs()` take a tuple signature and scans the `Rel.tsig` file, comparing that signature to the stored tuple signatures. It builds a bit-string showing which pages contain at least one "matching" tuple. It behaves roughly as follows:

```
QuerySig = makeTupleSig(Query)
Pages = AllZeroBits
foreach Tsig in tsigFile {
    if (Tsig matches QuerySig) {
        PID = data page for tuple corresponding to Tsig
        include PID in Pages
    }
}
```

Note that the i^{th} tuple in the data file has its corresponding signature as the i^{th} signature in the `Rel.tsig` file. However, since tuples and tuple signatures are different sizes, the page that the signature appears on will not necessarily have the same page ID as the page in which the corresponding tuple is located.

Task 4: Page Signatures (3 marks)

Implement indexing using page-level signatures (`psigs`).

This is similar to how tuple-level signature indexing is done, except that the signatures are larger. The functions that you need to complete are `makePageSig()` and `findPagesUsingPageSigs()` in the `psig.c` file. You will also need to add more code to the `addToRelation()` function to maintain page signatures when new tuples are inserted.

One major difference between tuple signatures and page signatures is that page signatures are not a one-off insertion. When a new tuple is added, its page-level signature needs to be included page signature for the page where it is inserted. The process can be described roughly as follows:

```
new Tuple is inserted into page PID
Psig = makePageSig(Tuple)
PPsig = fetch page signature for data page PID from psigFile
merge Psig and PPsig giving a new PPsig
update page signature for data page PID in psigFile
```

The `makePageSig()` function be used to generate a page-level signature for the query, and then used to generate a bit-string of matching pages roughly as follows:

```
QuerySig = makePageSig(Query)
Pages = AllZeroBits
foreach Psig in psigFile {
    if (Psig matches QuerySig) {
        PID = data page corresponding to Psig
        include PID in Pages
    }
}
```

Task 5: Bit-sliced Signatures (4 marks)

Implement indexing using bit-sliced page signatures.

Each bit-slice is effectively a list of pages that have a specific bit from the page-signature set to 1 (e.g. if a page-level signature has bit 5 set to one, the bit-slice 5 has a 1 bit for every page with a page signature where bit 5 is set). This drives both the updating of bit-slices and their use in indexing.

You will need to modify the functions: `newRelation()` in `reln.c`, `addToRelation()` in `reln.c`, and `findPagesUsingBitSlices()` in `bsig.c`. The modifications to `newRelation()` are relatively straightforward, but remember to update the relation parameters appropriately.

The `addToRelation()` should take a tuple, produce a page signature for it, then update all of the bit-slices corresponding to 1-bits in the page signature. This can be

```
PID = data page where
Psig = makePageSig(Tu
for each i in 0..pm-1 {
    if (Psig bit[i] is 1) {
        Slice = get i'th bit slice from bsigFile
        set the PID'th bit in Slice
        write updated Slice back to bsigfile
    }
}
```

The `findPagesUsingBitSlices()` function computes a page-level signature for query and then takes an intersection of the bit-slices corresponding to the 1-bits in the page signature. This gives a "matching" pages list straight away, and hopefully after reading far less of the `Rel.bsig` file than would be read using a `Rel.psig` file. The method can be described roughly as follows:

```
Qsig = makePageSig(Query)
Pages = AllOneBits
for each i in 0..pm-1 {
    if (Psig bit[i] is 1) {
        Slice = get i'th bit slice from bsigFile
        zero bits in Pages which are zero in Slice
    }
}
```

Task 6: Comparison of SIMC Indexing Methods (1 bonus mark)

For a range of parameter settings:

- size of tuples (determined by number of attributes)
- size of bit-slices (determined by expected number of tuples)
- size of tuple and page signatures (determined by p_F)

run a range of queries using the `select` command:

- not using indexing (use `x` as the signature type)
- open query (`?, ?, ..., ?`)
- query with one solution (`100001, ?, ..., ?`)
- query with many solutions (`?, ?, a3-001, ..., ?`)
- query with multiple values (`100001, ?, a3-001, ..., ?`)

Report the total number of page reads (signature pages + data pages), for each combination and make some comments about the effectiveness of each type of indexing under different scenarios.

Put your results in a file called `analysis.pdf` and submit that with your code files in the `ass2.tar` file.

Submission

You need to submit a single `tar` file containing all of the code files that are needed to build the `create`, `insert` and `select` commands, including a new `Makefile` if you add extra modules.

You should not submit the code for the commands, i.e. do not submit `create.c`, `insert.c` or `select.c`. We will use the original versions of these for testing your code, so you must preserve the module interfaces that these commands use.

When you want to submit your work, do the following:

```
$ cd your/ass2/directory
$ tar cf ass2.tar FilesToSubmit
```

The `FilesToSubmit` would typically include:

```
Makefile
bits.c bits.h page.c page.h query.c query.h
reln.c reln.h tuple.c tuple.h util.c util.h
tsig.c tsig.h psig.c psig.c bsig.c bsig.h
```

plus any other modules you may have created, along with the modified `Makefile` to include them in the compilation.

Once you have generated the `ass2.tar` file, you can submit it via WebCMS.

Have fun, *jas*

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat `edu_assist_pro`